

SMARTBRIDGE EXTERNSHIP - APPLIED DATA SCIENCE

ASSIGNMENT-3

Name: Peddibhotla Sree Rakshitha
Registration Number: 20BCD7095
Campus: VIT-AP University

ADS Assignment 3

Problem Statement: House Price Prediction

Description:- House price prediction is a common problem in the real estate industry and involves predicting the selling price of a house based on various features and attributes. The problem is typically approached as a regression problem, where the target variable is the price of the house, and the features are various attributes of the house

The features used in house price prediction can include both quantitative and categorical variables, such as the number of bedrooms, house area, bedrooms, furnished, nearness to main road, and various amenities such as a garage and other factors that may influence the value of the property.

Accurate predictions can help agents and appraisers price homes correctly, while homeowners can use the predictions to set a reasonable asking price for their properties. Accurate house price prediction can also be useful for buyers who are looking to make informed decisions about purchasing a property and obtaining a fair price for their investment.

Attribute Information:

- Name - Description
- 1- Price-Prices of the houses
- 2- Area- Area of the houses
- 3- Bedrooms- No of house bedrooms
- 4- Bathrooms- No of bathrooms
- 5- Stories- No of house stories
- 6- Main Road- Weather connected to Main road
- 7- Guestroom-Weather has a guest room
- 8- Basement-Weather has a basement
- 9- Hot water heating- Weather has a hot water heater
- 10-Airconditioning-Weather has a air conditioner
- 11-Parking- No of house parking
- 12-Furnishing Status-Furnishing status of house

Building a Regression Model

1. Download the dataset: Dataset

2. Load the dataset into the tool.

```
✓ [1] import pandas as pd
    import numpy as np

✓ [2] df=pd.read_csv('Housing.csv')

✓ [3] df
   ▾
   price area bedrooms bathrooms stories mainroad guestroom basement hotwaterheating airconditioning parking furnishingstatus ✎
   0 13300000 7420 4 2 3 yes no no no yes 2 furnished
   1 12250000 8960 4 4 4 yes no no no yes 3 furnished
   2 12250000 9960 3 2 2 yes no yes no no 2 semi-furnished
   3 12215000 7500 4 2 2 yes no yes no yes 3 furnished
   4 11410000 7420 4 1 2 yes yes yes no yes 2 furnished
   ...
   540 1820000 3000 2 1 1 yes no yes no no 2 unfurnished
   541 1767150 2400 3 1 1 no no no no no 0 semi-furnished
   542 1750000 3620 2 1 1 yes no no no no 0 unfurnished
   543 1750000 2910 3 1 1 no no no no no 0 furnished
   544 1750000 3850 3 1 2 yes no no no no 0 unfurnished
545 rows × 12 columns

✓ [4] df.dtypes
   ▾
   price int64
   area int64
   bedrooms int64
   bathrooms int64
   stories int64
   mainroad object
   guestroom object
   basement object
   hotwaterheating object
   airconditioning object
   parking int64
   furnishingstatus object
   dtype: object

✓ [5] df.head()
   ▾
   price area bedrooms bathrooms stories mainroad guestroom basement hotwaterheating airconditioning parking furnishingstatus ✎
   0 13300000 7420 4 2 3 yes no no no yes 2 furnished
   1 12250000 8960 4 4 4 yes no no no yes 3 furnished
   2 12250000 9960 3 2 2 yes no yes no no 2 semi-furnished
   3 12215000 7500 4 2 2 yes no yes no yes 3 furnished
   4 11410000 7420 4 1 2 yes yes yes no yes 2 furnished

✓ [6] df.replace('?',np.nan,inplace=True)

✓ [7] df.replace('?',np.nan,inplace=True)

✓ [8] df.shape
(545, 12)

✓ [9] df.describe()
   ▾
   price area bedrooms bathrooms stories parking ✎
   count 5.450000e+02 545.000000 545.000000 545.000000 545.000000
   mean 4.766729e+06 5150.541284 2.965138 1.286239 1.805505 0.693578
   std 1.870440e+06 2170.141023 0.738064 0.502470 0.867492 0.861586
   min 1.750000e+06 1650.000000 1.000000 1.000000 1.000000 0.000000
   25% 3.430000e+06 3600.000000 2.000000 1.000000 1.000000 0.000000
   50% 4.340000e+06 4600.000000 3.000000 1.000000 2.000000 0.000000
   75% 5.740000e+06 6360.000000 3.000000 2.000000 2.000000 1.000000
   max 1.330000e+07 16200.000000 6.000000 4.000000 4.000000 3.000000

✓ [10] df.info()
   ▾
   <class 'pandas.core.frame.DataFrame'>
   RangeIndex: 545 entries, 0 to 544
   Data columns (total 12 columns):
```

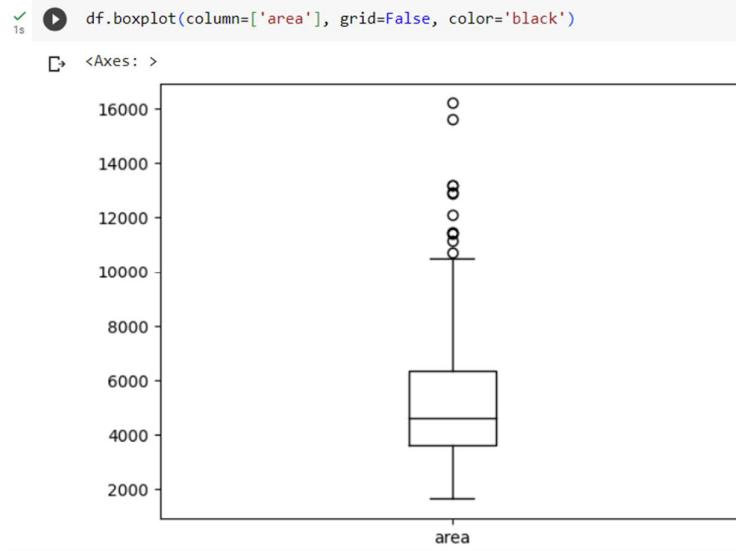
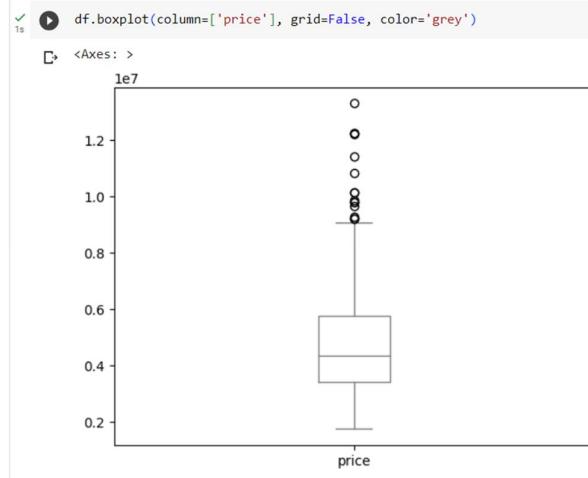
```

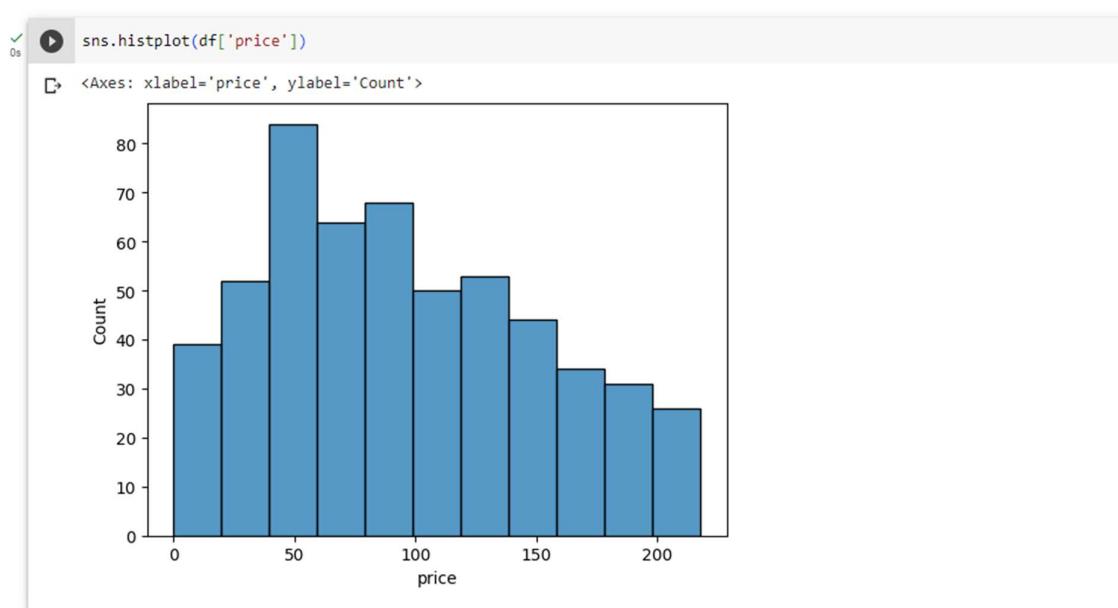
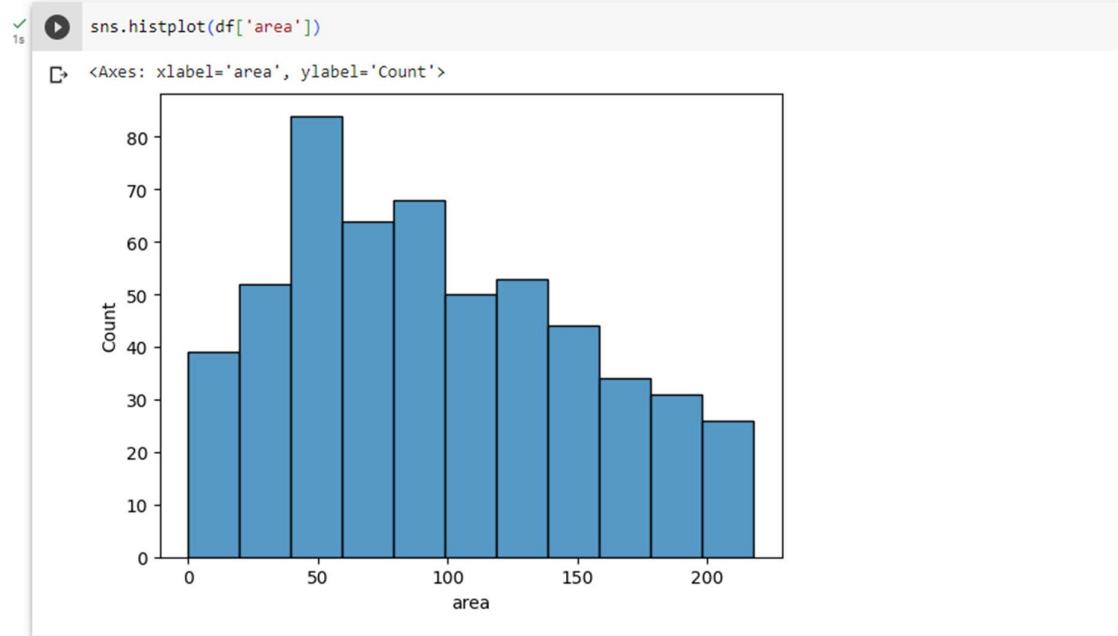
0s df.info()
0s <class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   price       545 non-null    int64  
 1   area        545 non-null    int64  
 2   bedrooms    545 non-null    int64  
 3   bathrooms   545 non-null    int64  
 4   stories     545 non-null    int64  
 5   mainroad    545 non-null    object  
 6   guestroom   545 non-null    object  
 7   basement    545 non-null    object  
 8   hotwaterheating 545 non-null    object  
 9   airconditioning 545 non-null    object  
 10  parking     545 non-null    int64  
 11  furnishingstatus 545 non-null    object  
dtypes: int64(6), object(6)
memory usage: 51.2+ KB

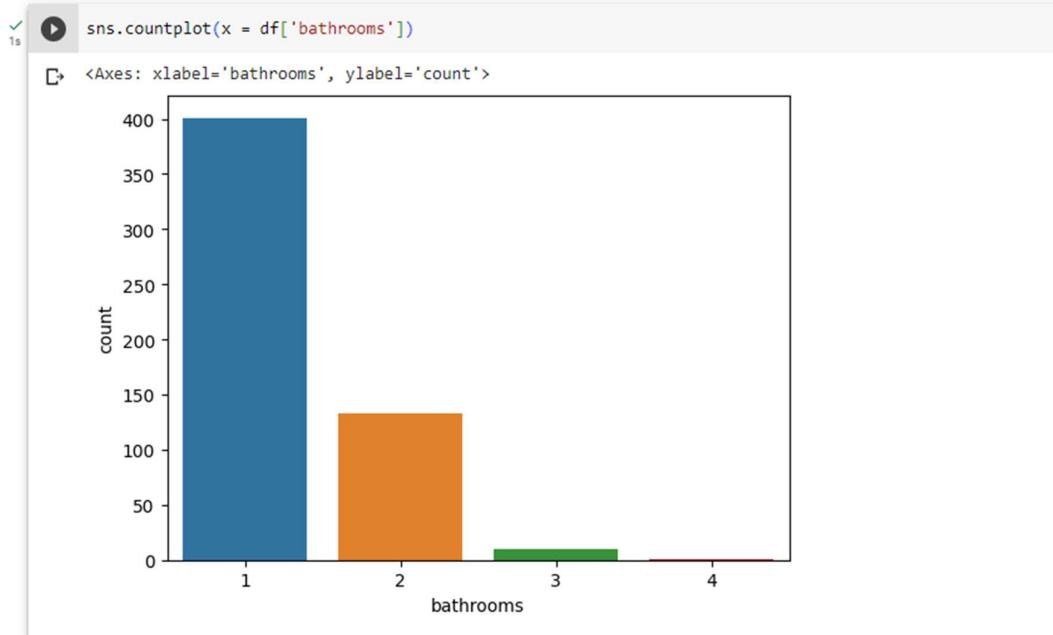
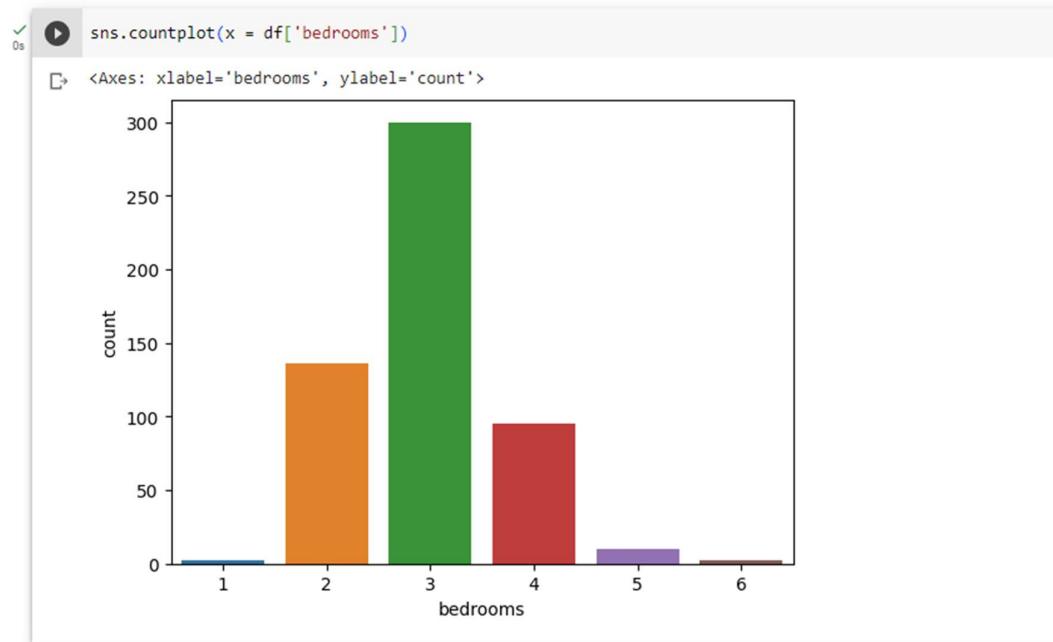
```

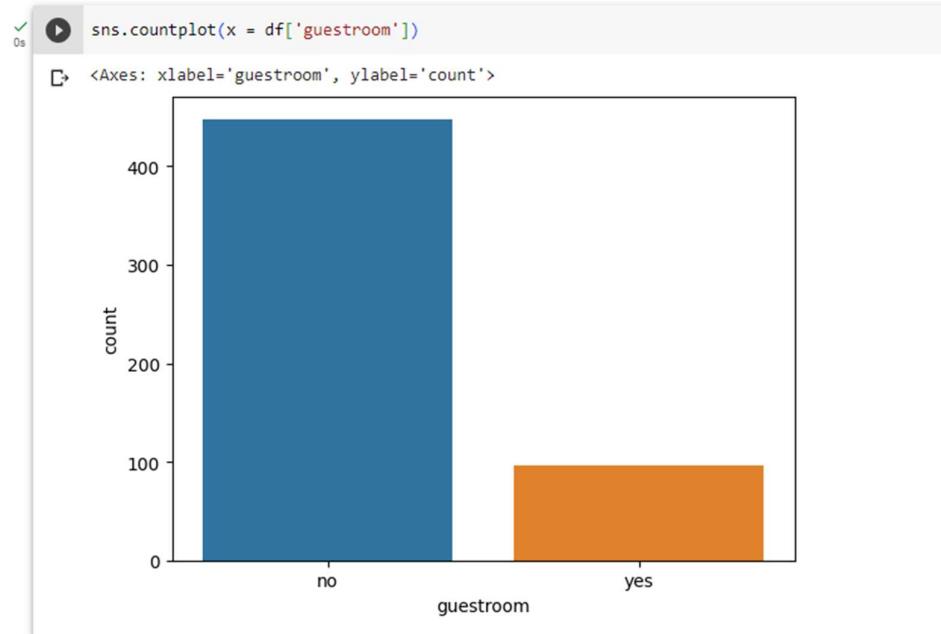
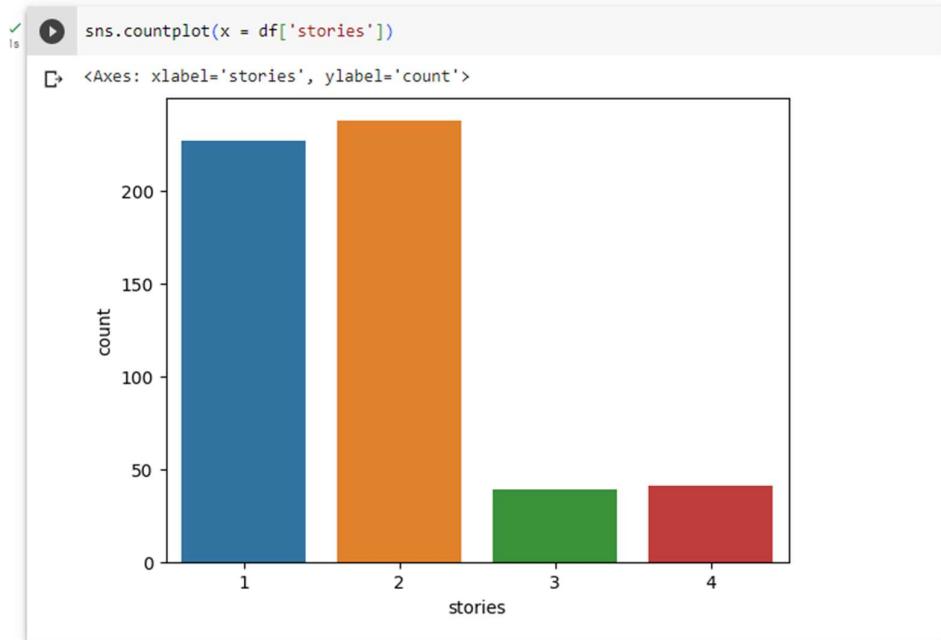
3. Perform Below Visualizations.

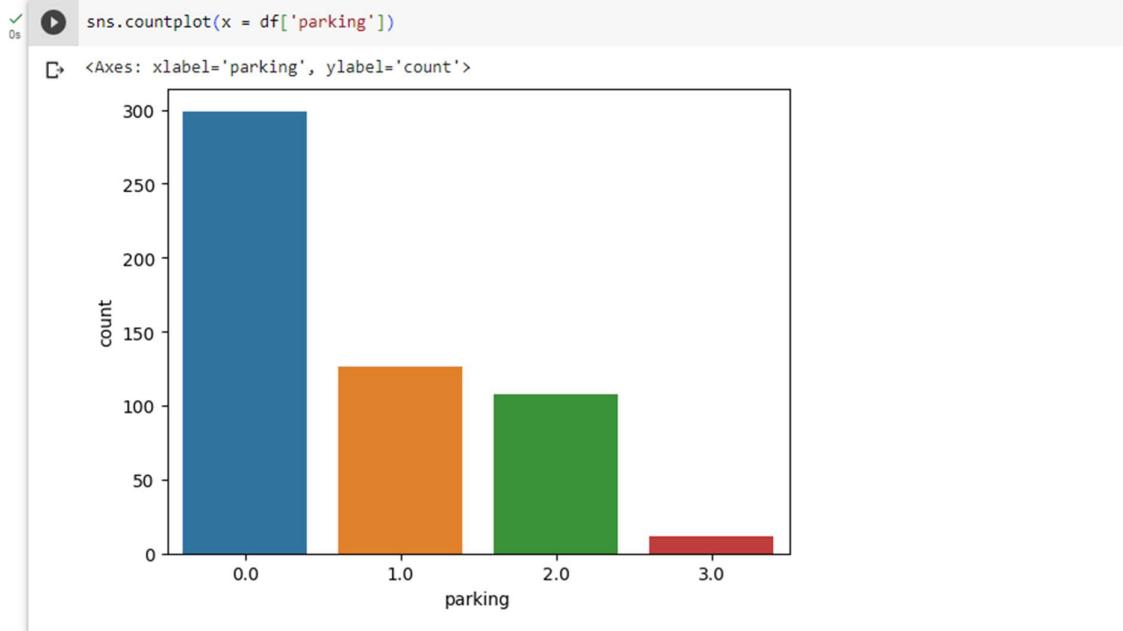
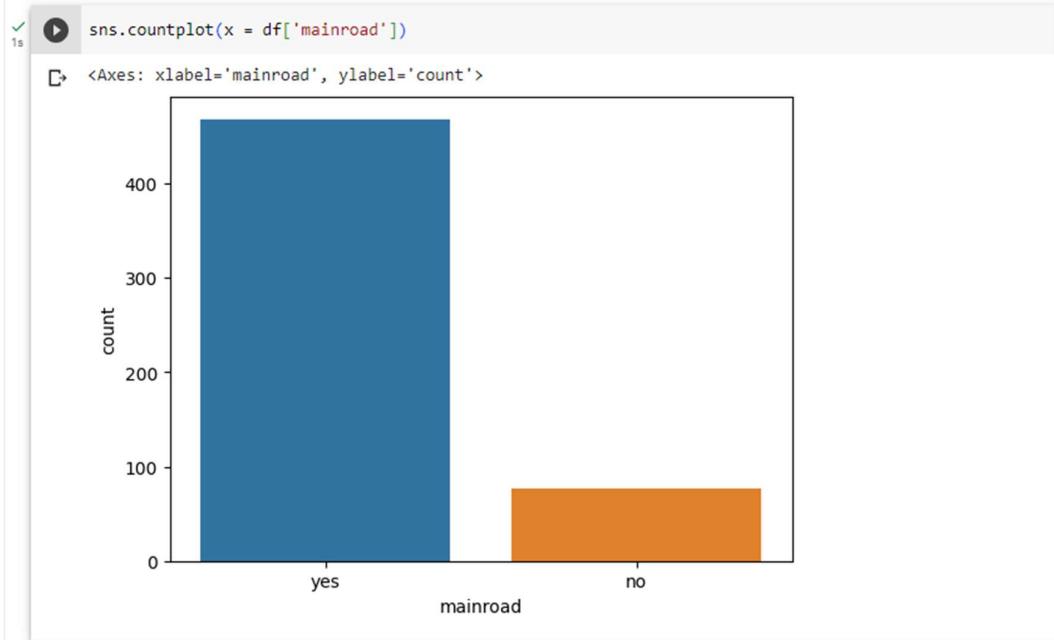
- Univariate Analysis

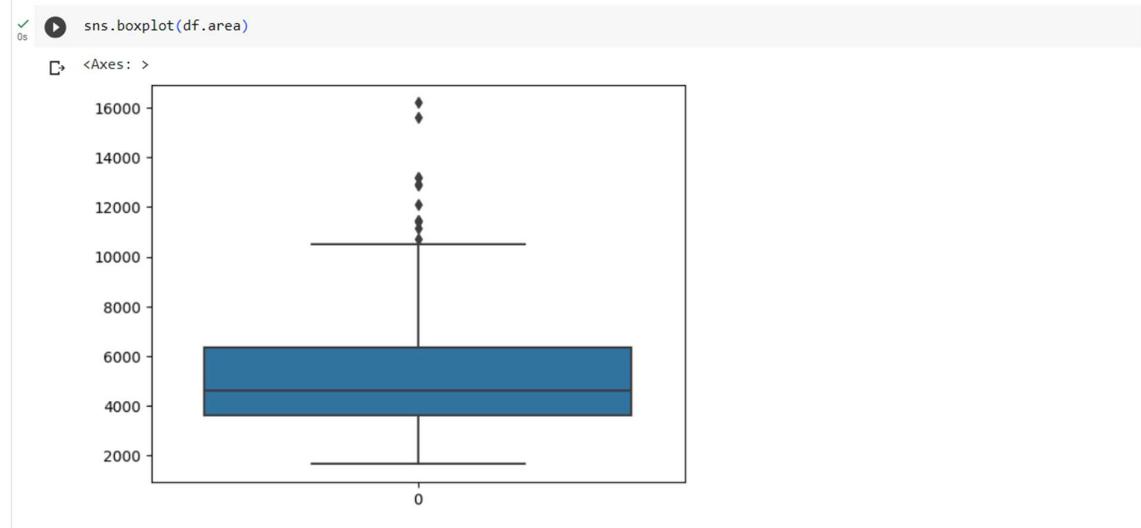
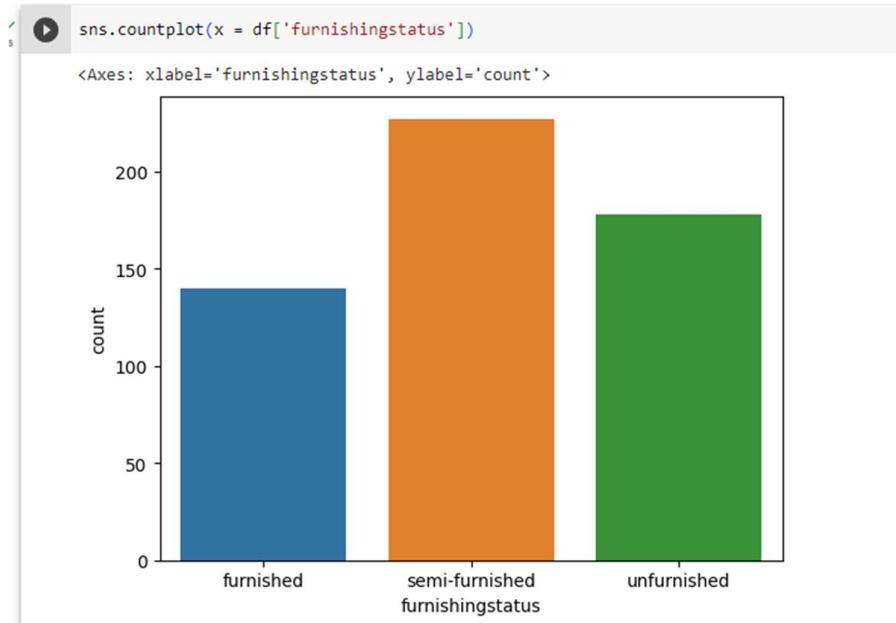


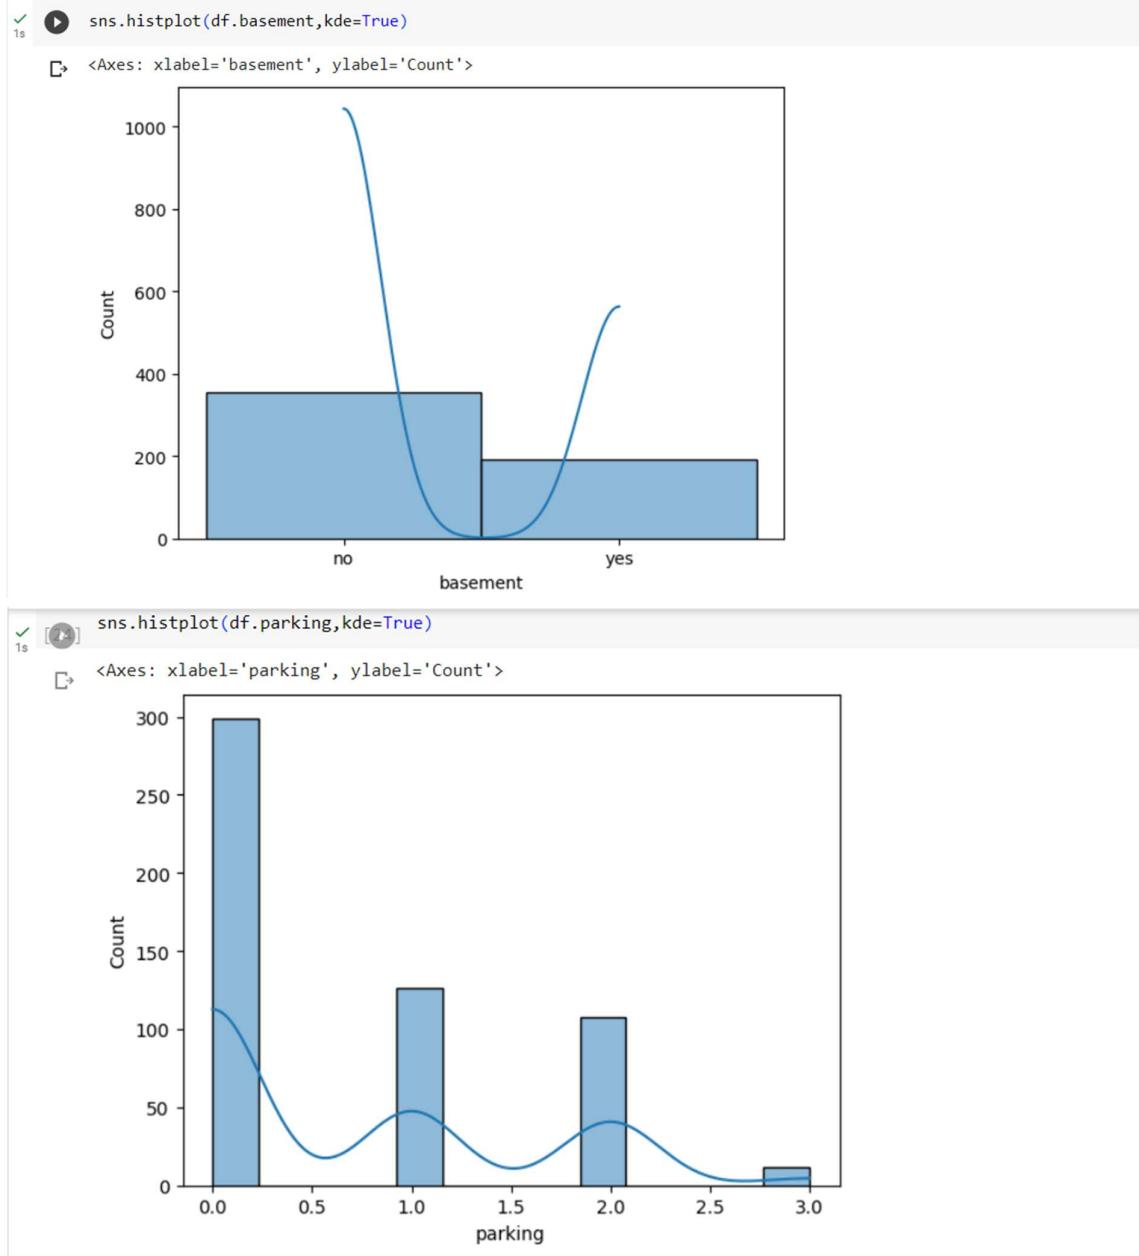






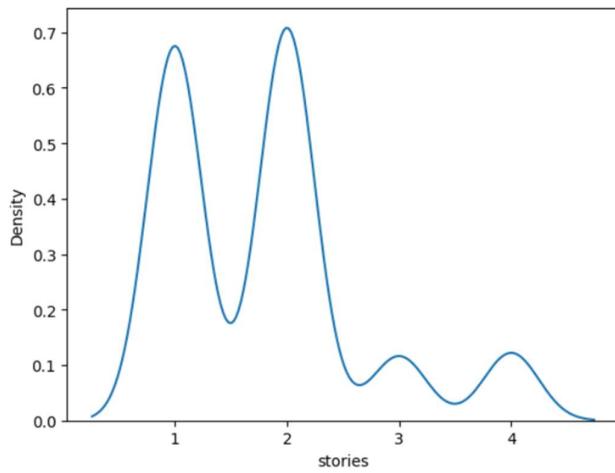






```
✓ 0s   sns.kdeplot(df['stories'])
```

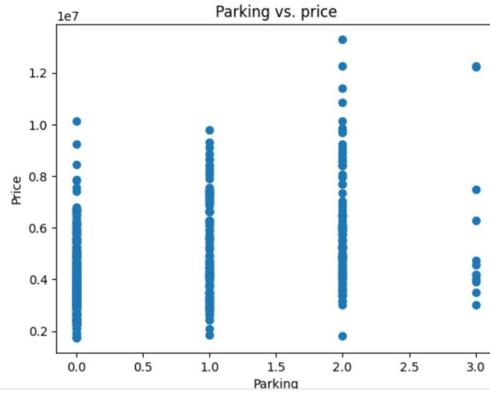
```
↳ <Axes: xlabel='stories', ylabel='Density'>
```

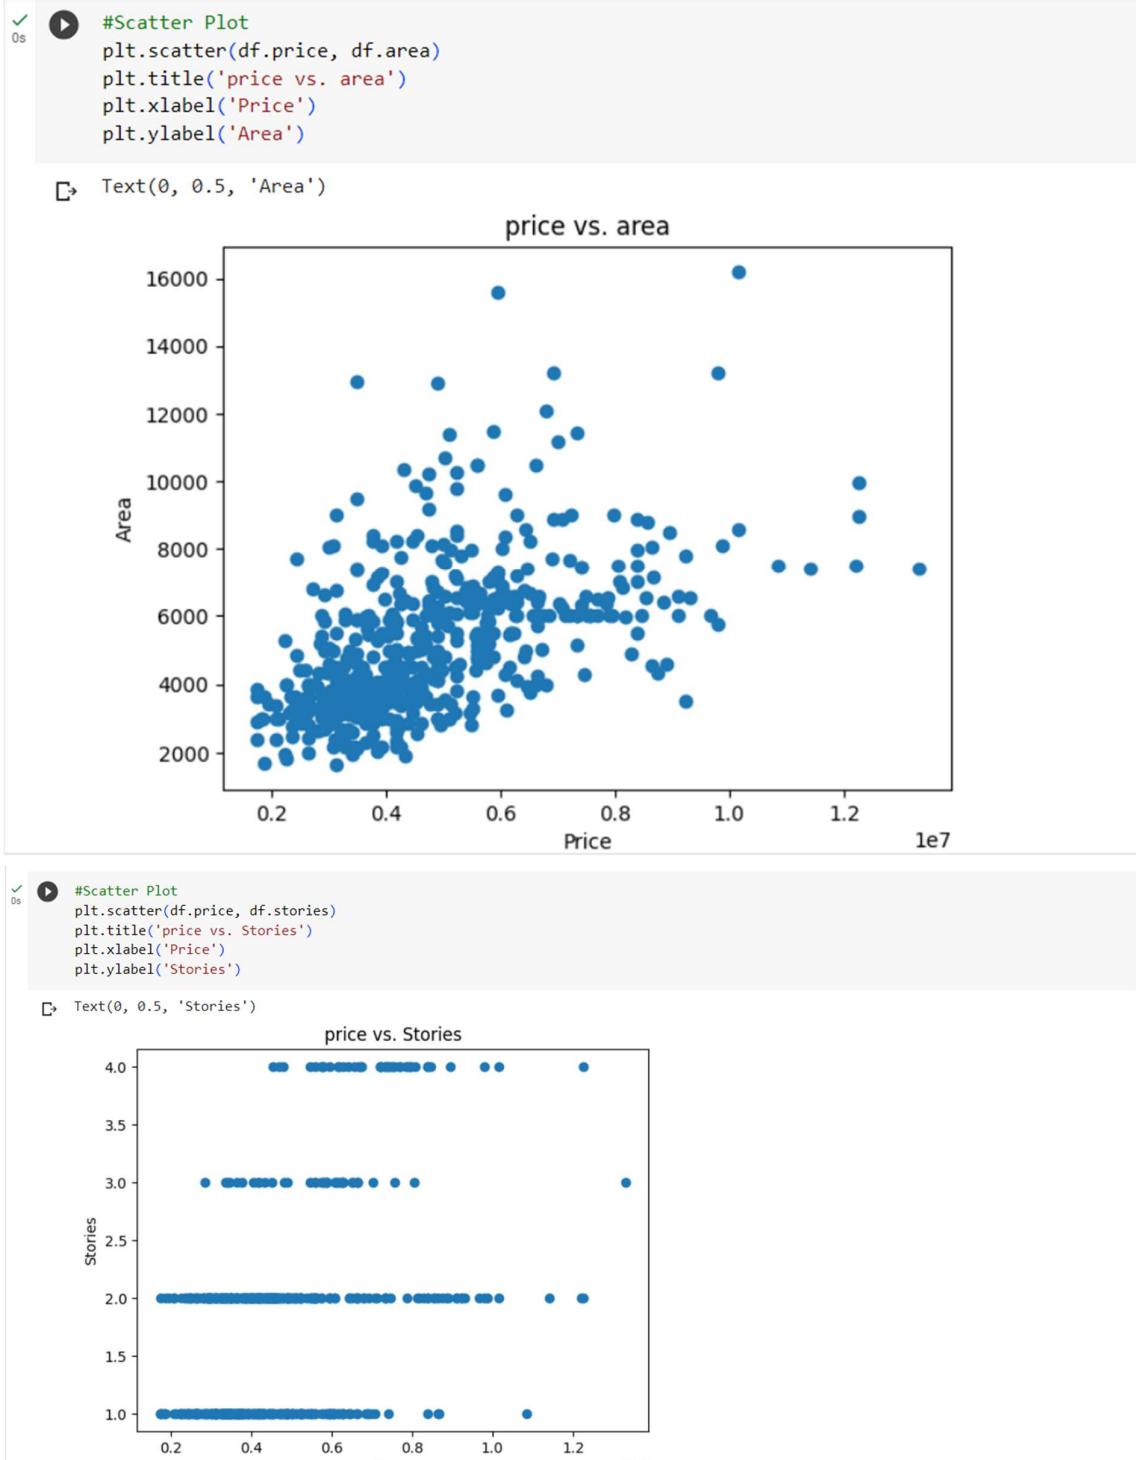


- Bi-Variate Analysis

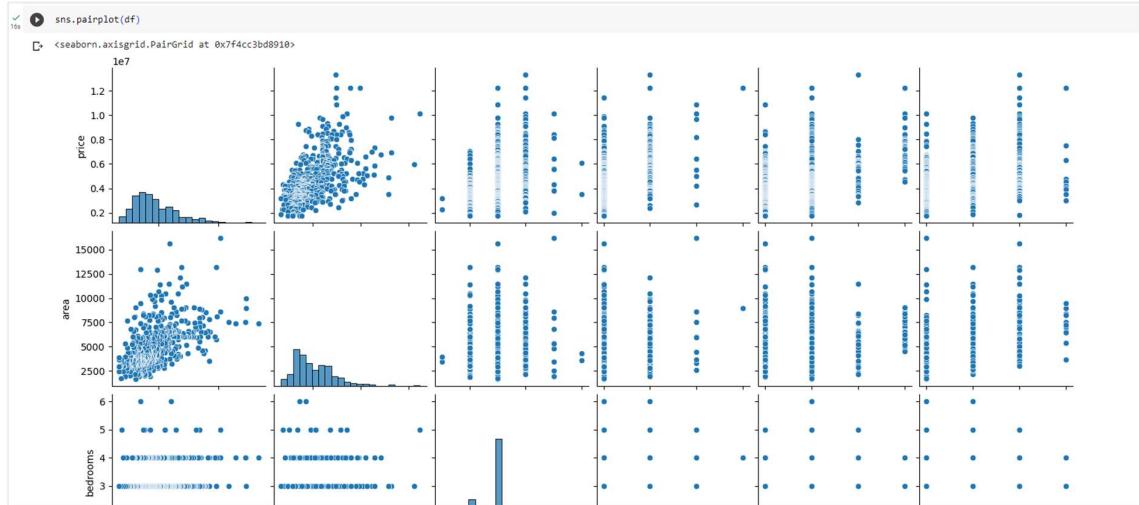
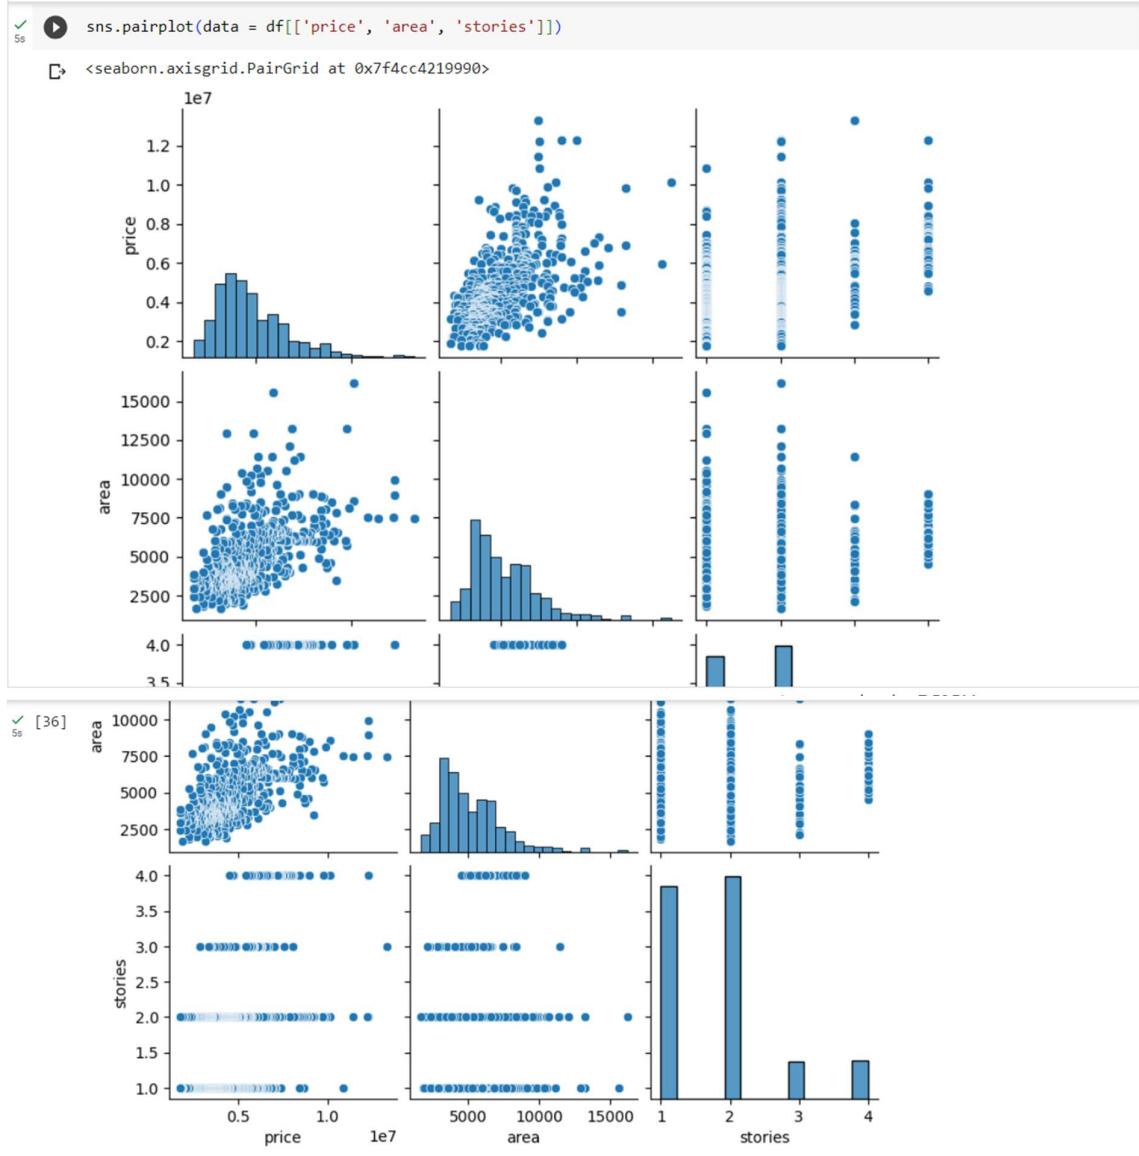
```
✓ 0s   #Scatter Plot  
      plt.scatter(df.parking, df.price)  
      plt.title('Parking vs. price')  
      plt.xlabel('Parking')  
      plt.ylabel('Price')
```

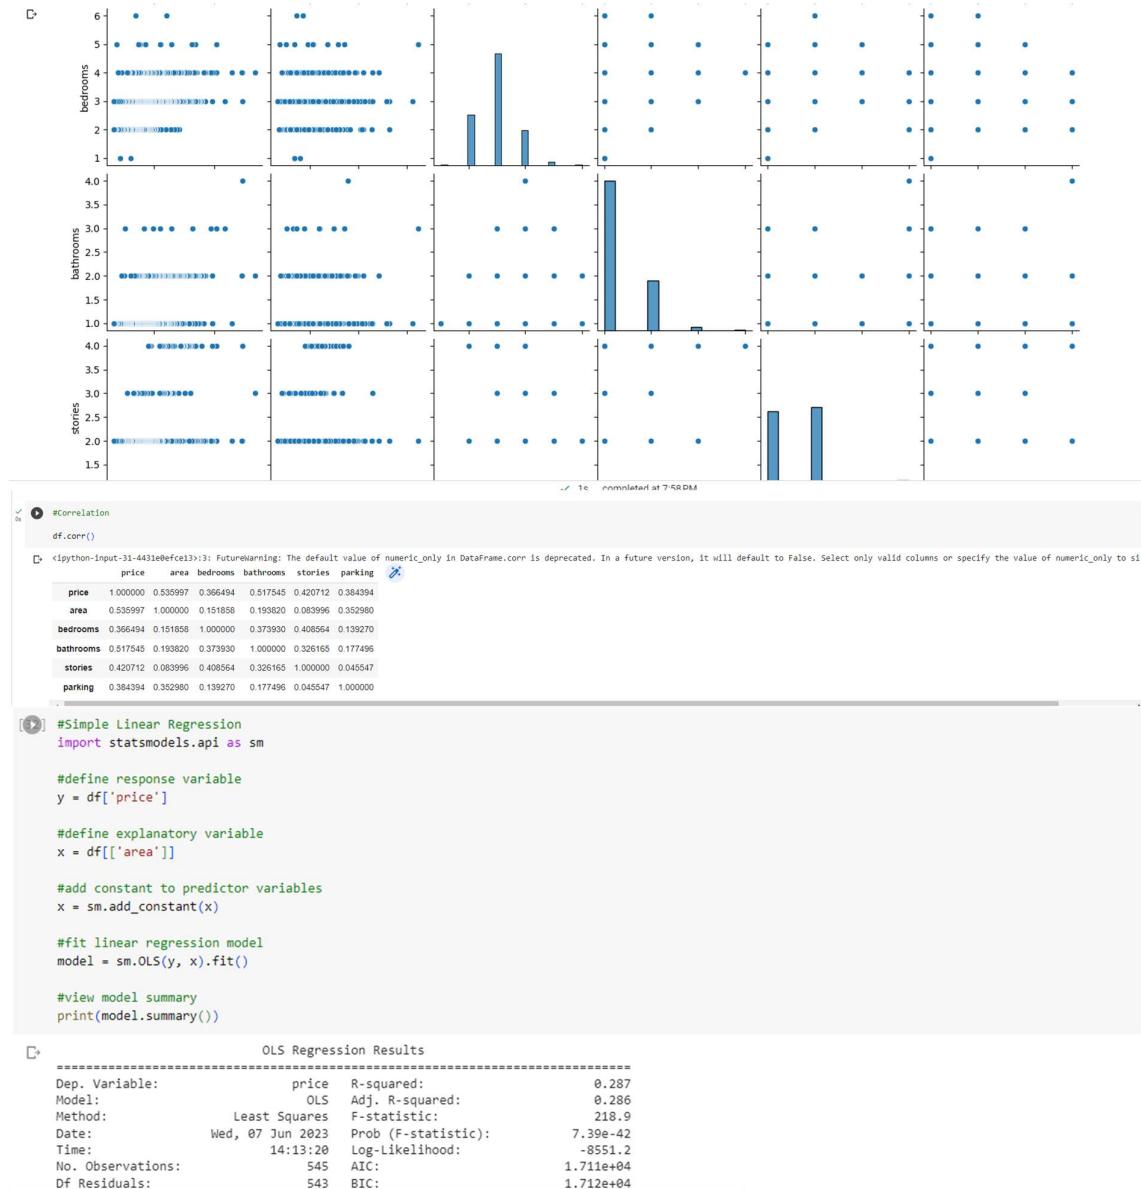
```
↳ Text(0, 0.5, 'Price')
```





- Multi-Variate Analysis





4. Perform descriptive statistics on the dataset.

```

✓ [40] stats = df.describe(include = 'all')
print(stats)

   price      area  bedrooms  bathrooms  stories \
count  5.450000e+02  545.000000  545.000000  545.000000
unique       NaN      NaN      NaN      NaN      NaN
top        NaN      NaN      NaN      NaN      NaN
freq       NaN      NaN      NaN      NaN      NaN
mean    4.766720e+06  5150.541284  2.965138  1.286239  1.805505
std     1.870440e+06  2170.141023  0.738064  0.582470  0.867492
min    1.750000e+06  1650.000000  1.000000  1.000000  1.000000
25%   3.430000e+06  3600.000000  2.000000  1.000000  1.000000
50%   4.340000e+06  4600.000000  3.000000  1.000000  2.000000
75%   5.740000e+06  6360.000000  3.000000  2.000000  2.000000
max   1.330000e+07  16200.000000  6.000000  4.000000  4.000000

   mainroad  guestroom  basement  hotwaterheating  airconditioning \
count      545          545          545          545          545
unique       2            2            2            2            2
top        yes           no           no           no           no
freq       468          448          354          520          373
mean      NaN          NaN          NaN          NaN          NaN
std       NaN          NaN          NaN          NaN          NaN
min       NaN          NaN          NaN          NaN          NaN
25%      NaN          NaN          NaN          NaN          NaN
50%      NaN          NaN          NaN          NaN          NaN
75%      NaN          NaN          NaN          NaN          NaN
max       NaN          NaN          NaN          NaN          NaN

   parking  furnishingstatus
count  545.000000          545
unique       NaN                  3
top        NaN  semi-furnished
freq       NaN                  227
mean    0.693578          NaN
std     0.861586          NaN
min    0.000000          NaN
25%   0.000000          NaN
50%   0.000000          NaN
75%   1.000000          NaN
dtype: float64

✓ [40] df.mean()

<ipython-input-40-c61f0c8cf89b5>:1: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. I
price    4.766720e+06
area     5.150541e+03
bedrooms  2.965138e+00
bathrooms 1.286239e+00
stories   1.805505e+00
parking   6.935780e-01
dtype: float64

✓ [41] df.median()

   <ipython-input-41-6d467abf240d>:1: FutureWarning: The default value of numeric_only in DataFrame.median is deprecated. In
price      4340000.0
area       4600.0
bedrooms    3.0
bathrooms   1.0
stories      2.0
parking      0.0
dtype: float64

✓ [42] df.max()

   <ipython-input-42-6d467abf240d>:1: FutureWarning: The default value of numeric_only in DataFrame.max is deprecated. In
price      13300000.0
area       16200.0
bedrooms     6.0
bathrooms    4.0
stories      4.0
mainroad     yes
guestroom    yes
basement     yes
hotwaterheating  yes
airconditioning  yes
parking      3.0
furnishingstatus  unfurnished
dtype: object

✓ [43] df.std()

<ipython-input-43-ce97bb7eaef8>:1: FutureWarning: The default value of numeric_only in DataFrame.std is deprecated. In a future version,
price    1.870440e+06
area     2.170141e+03
bedrooms  7.380639e-01
bathrooms 5.024696e-01
stories   8.674925e-01
parking   8.615858e-01
dtype: float64

```

✓ [43] df.isnull()

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	furnishingstatus
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False
...
540	False	False	False	False	False	False	False	False	False	False	False	False
541	False	False	False	False	False	False	False	False	False	False	False	False
542	False	False	False	False	False	False	False	False	False	False	False	False
543	False	False	False	False	False	False	False	False	False	False	False	False
544	False	False	False	False	False	False	False	False	False	False	False	False

545 rows × 12 columns

5. Check for Missing values and deal with them.

✓ [44] df.isnull()

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	furnishingstatus
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False
...
540	False	False	False	False	False	False	False	False	False	False	False	False
541	False	False	False	False	False	False	False	False	False	False	False	False
542	False	False	False	False	False	False	False	False	False	False	False	False
543	False	False	False	False	False	False	False	False	False	False	False	False
544	False	False	False	False	False	False	False	False	False	False	False	False

545 rows × 12 columns

✓ [45] pd.isnull(df["parking"])

0	False
1	False
2	False
3	False
4	False
...	...

✓ [46] df['area'].fillna(df['area'].mean(), inplace=True)

6. Find the outliers and replace them outliers

✓ [47] from scipy import stats
import numpy as np
z = np.abs(stats.zscore(df['parking']))
print(z)

0	1.517692
1	2.679489
2	1.517692
3	2.679489
4	1.517692
...	...
540	1.517692
541	0.805741
542	0.805741
543	0.805741
544	0.805741

Name: parking, Length: 545, dtype: float64

✓ [48] threshold = 2

Position of the outlier
print(np.where(z > 2))

(array([1, 3, 47, 93, 225, 247, 299, 304, 323, 331, 401, 472]),)

✓ [49] Q1 = np.percentile(df['parking'], 25, method='midpoint')
Q3 = np.percentile(df['parking'], 75, method='midpoint')
IQR = Q3 - Q1
print(IQR)

1.0

```

✓ [50] # Above Upper bound
    upper=Q3+1.5*IQR
    upper_array=np.array(df['parking']>=upper)
    print("Upper Bound:",upper)
    print(upper_array.sum())

    #Below Lower bound
    lower=Q1-1.5*IQR
    lower_array=np.array(df['parking']<=lower)
    print("Lower Bound:",lower)
    print(lower_array.sum())

Upper Bound: 2.5
12
Lower Bound: -1.5
0

✓ 0s ⏪ print("Old Shape: ", df.shape)

    ''' Detection '''
    # IQR
    # Calculate the upper and lower limits
    Q1 = df['parking'].quantile(0.25)
    Q3 = df['parking'].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5*IQR
    upper = Q3 + 1.5*IQR

    # Create arrays of Boolean values indicating the outlier rows
    upper_array = np.where(df['parking']>=upper)[0]
    lower_array = np.where(df['parking']<=lower)[0]

    # Removing the outliers
    df['parking'].drop(index=upper_array, inplace=True)
    df['parking'].drop(index=lower_array, inplace=True)

    # Create arrays of Boolean values indicating the outlier rows
    upper_array = np.where(df['parking']>=upper)[0]
    lower_array = np.where(df['parking']<=lower)[0]

    # Removing the outliers
    df['parking'].drop(index=upper_array, inplace=True)
    df['parking'].drop(index=lower_array, inplace=True)

    # Print the new shape of the DataFrame
    print("New Shape: ", df['parking'].shape)

Old Shape: (545, 12)
New Shape: (545,)

✓ 0s [53] median = df.loc[df['parking']<13, 'parking'].median()
    df.loc[df.parking > 13, 'parking'] = np.nan
    df.fillna(median,inplace=True)

```

7. Check for Categorical columns and perform encoding.

```

✓ 0s ⏪ df_main.corr()

D: <ipython-input-61-b764c75a6398>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns.
      price     area bedrooms bathrooms parking stories_1 stories_2 stories_3 stories_4
price  1.000000  1.000000  0.373302  0.493090  0.363694 -0.266898 -0.017796  0.158267  0.377645
area   1.000000  1.000000  0.373302  0.493090  0.363694 -0.266898 -0.017796  0.158267  0.377645
bedrooms 0.373302  0.373302  1.000000  0.373930  0.139270 -0.510235  0.377769  0.100010  0.145561
bathrooms 0.493090  0.493090  0.373930  1.000000  0.177496 -0.251906  0.080144  0.040225  0.280809
parking  0.363694  0.363694  0.139270  0.177496  1.000000  0.015383 -0.043285 -0.074838  0.125775
stories_1 -0.266898 -0.266898 -0.510235 -0.251905  0.015383  1.000000 -0.743907 -0.234561 -0.240977
stories_2 -0.017796 -0.017796  0.377769  0.080144 -0.043285 -0.743907  1.000000 -0.244442 -0.251128
stories_3  0.158267  0.158267  0.100010  0.040225 -0.074838 -0.234561 -0.244442  1.000000 -0.079183
stories_4  0.377645  0.377645  0.145561  0.280809  0.125775 -0.240977 -0.251128 -0.079183  1.000000

```

```

X
✓ [55] #label encoding
  from sklearn.preprocessing import LabelEncoder
  le=LabelEncoder()
  df.price=le.fit_transform(df.price)

✓ [57] df.area=df.fit_transform(df.price)

✓ [58] df.head()

  price area bedrooms bathrooms stories mainroad guestroom basement hotwaterheating airconditioning parking furnishingstatus
  0 218 218 4 2 3 yes no no no yes 2.0 furnished
  1 217 217 4 4 4 yes no no no yes 3.0 furnished
  2 217 217 3 2 2 yes no yes no no 2.0 semi-furnished
  3 216 216 4 2 2 yes no yes no yes 3.0 furnished
  4 215 215 4 1 2 yes yes yes no yes 2.0 furnished

✓ [59] df_main=pd.get_dummies(df,columns=['stories'])

✓ [60] df_main.head()

  price area bedrooms bathrooms mainroad guestroom basement hotwaterheating airconditioning parking furnishingstatus stories_1 stories_2 stories_3 stories_4
  0 218 218 4 2 yes no no no yes 2.0 furnished 0 0 1 0
  1 217 217 4 4 yes no no no yes 3.0 furnished 0 0 0 1
  2 217 217 3 2 yes no yes no no 2.0 semi-furnished 0 1 0 0
  3 216 216 4 2 yes no yes no yes 3.0 furnished 0 1 0 0
  4 215 215 4 1 yes yes yes no yes 2.0 furnished 0 1 0 0

✓ [61] #Heatmap
  plt.figure(figsize=(10,8))
  sns.heatmap(df_main.corr(),annot=True)

  ▶ - 1 1 0.37 0.49 0.36 -0.27 -0.018 0.16 0.38
  ▶ - 1 1 0.37 0.49 0.36 -0.27 -0.018 0.16 0.38
  ▶ - 0.37 0.37 1 0.37 0.14 -0.51 0.38 0.1 0.15
  ▶ - 0.49 0.49 0.37 1 0.18 -0.25 0.08 0.04 0.28
  ▶ - 0.36 0.36 0.14 0.18 1 0.015 -0.043 -0.075 0.13
  ▶ - -0.27 -0.27 -0.51 -0.25 0.015 1 -0.74 -0.23 -0.24
  ▶ - -0.018 -0.018 0.38 0.08 -0.043 -0.74 1 -0.24 -0.25
  ▶ - 0.16 0.16 0.1 0.04 -0.075 -0.23 -0.24 1 -0.079
  ▶ - 0.38 0.38 0.15 0.28 0.13 -0.24 -0.25 -0.079 1

  price area bedrooms bathrooms mainroad guestroom basement hotwaterheating airconditioning parking furnishingstatus stories_1 stories_2 stories_3 stories_4

```

8. Split the data
into dependent and independent variables.

```

✓ [63] # Y is dependent variable(target)
0s y=df_main['price']
y.head()

0    218
1    217
2    217
3    216
4    215
Name: price, dtype: int64

✓ [65] # X is independent variable or predictors
0s X=df_main.drop(columns=['price', 'hotwaterheating', 'basement', 'airconditioning', 'furnishingstatus', 'mainroad', 'guestroom'],axis=1)
X.head()

   area  bedrooms  bathrooms  parking  stories_1  stories_2  stories_3  stories_4
0    218        4         2      2.0        0        0        1        0
1    217        4         4      3.0        0        0        0        1
2    217        3         2      2.0        0        1        0        0
3    216        4         2      3.0        0        1        0        0
4    215        4         1      2.0        0        1        0        0

```

9. Scale the independent variables

```

✓ [66] name=X.columns
0s name

Index(['area', 'bedrooms', 'bathrooms', 'parking', 'stories_1', 'stories_2',
       'stories_3', 'stories_4'],
      dtype='object')

✓ [67] from sklearn.preprocessing import MinMaxScaler
0s scale=MinMaxScaler()

✓ [68] X_scaled=scale.fit_transform(X)
0s X_scaled

array([[1.          , 0.6        , 0.33333333, ..., 0.          ,
       1.          , 0.          ],
       [0.99541284, 0.6        , 1.          , ..., 0.          ,
       0.          , 0.          ],
       [0.99541284, 0.4        , 0.33333333, ..., 1.          ,
       0.          , 0.          ],
       ...,
       [0.          , 0.2        , 0.          , ..., 0.          ,
       0.          , 0.          ],
       [0.          , 0.4        , 0.          , ..., 0.          ,
       0.          , 0.          ],
       [0.          , 0.4        , 0.          , ..., 1.          ,
       0.          , 0.          ],
       [0.          , 0.          , 0.          , ..., 0.          ,
       0.          , 0.          ]])

✓ [69] X=pd.DataFrame(X_scaled,columns=name)
0s X

   area  bedrooms  bathrooms  parking  stories_1  stories_2  stories_3  stories_4
0    1.000000    0.6    0.333333  0.666667        0.0        0.0        1.0        0.0
1    0.995413    0.6    1.000000  1.000000        0.0        0.0        0.0        1.0
2    0.995413    0.4    0.333333  0.666667        0.0        1.0        0.0        0.0

```

```

[69] X=pd.DataFrame(X_scaled,columns=name)
      X

      area bedrooms bathrooms parking stories_1 stories_2 stories_3 stories_4
0 1.000000 0.6 0.333333 0.666667 0.0 0.0 1.0 0.0
1 0.995413 0.6 1.000000 1.000000 0.0 0.0 0.0 1.0
2 0.995413 0.4 0.333333 0.666667 0.0 1.0 0.0 0.0
3 0.990826 0.6 0.333333 1.000000 0.0 1.0 0.0 0.0
4 0.986239 0.6 0.000000 0.666667 0.0 1.0 0.0 0.0
...
540 0.009174 0.2 0.000000 0.666667 1.0 0.0 0.0 0.0
541 0.004587 0.4 0.000000 0.000000 1.0 0.0 0.0 0.0
542 0.000000 0.2 0.000000 0.000000 1.0 0.0 0.0 0.0
543 0.000000 0.4 0.000000 0.000000 1.0 0.0 0.0 0.0
544 0.000000 0.4 0.000000 0.000000 0.0 1.0 0.0 0.0

```

545 rows × 8 columns

10. Split the data into training and testing

```

[70] from sklearn.model_selection import train_test_split
[71] X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
[72] X_train.head()

      area bedrooms bathrooms parking stories_1 stories_2 stories_3 stories_4
542 0.000000 0.2 0.0 0.0 1.0 0.0 0.0 0.0
496 0.100917 0.2 0.0 0.0 1.0 0.0 0.0 0.0
484 0.123853 0.2 0.0 0.0 1.0 0.0 0.0 0.0
507 0.082569 0.2 0.0 0.0 1.0 0.0 0.0 0.0
252 0.435780 0.4 0.0 0.0 1.0 0.0 0.0 0.0

[73] X_test.head()

      area bedrooms bathrooms parking stories_1 stories_2 stories_3 stories_4
239 0.449541 0.4 0.0 0.333333 0.0 1.0 0.0 0.0
113 0.674312 0.4 0.0 0.666667 1.0 0.0 0.0 0.0
325 0.330275 0.6 0.0 0.000000 0.0 1.0 0.0 0.0
66 0.788991 0.2 0.0 0.333333 1.0 0.0 0.0 0.0
479 0.128440 0.6 0.0 0.000000 0.0 1.0 0.0 0.0

y_train
542 0
496 22
484 27

```

```
✓ [74] y_train
0s
 542      0
496     22
484     27
507     18
252    95
...
70    169
277    85
9    211
359    62
192   117
Name: price, Length: 436, dtype: int64
```

```
✓ [75] y_test
0s
 239    98
113   147
325    72
66    172
479    28
...
76    165
132   140
311    77
464    34
155   133
Name: price, Length: 109, dtype: int64
```

-
11. Build the Model
 12. Train the Model
 13. Test the Model
 14. Measure the performance using Metrics.

```
✓ [76] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=101)

✓ [77] from sklearn.model_selection import train_test_split
       from sklearn.linear_model import LinearRegression
       from sklearn import metrics
       %matplotlib inline

✓ [78] lm = LinearRegression()

✓ [79] lm.fit(X_train,y_train)
       + LinearRegression()
       LinearRegression()

✓ [80] print(lm.intercept_)
2.842170943040401e-14

✓ [81] lm.coef_
array([ 2.18000000e+02,  5.09309693e-14,  6.56039411e-14, -9.06361449e-14,
       -3.53024649e-14, -8.03981922e-15, -1.75456812e-14, -4.83003205e-14])
```

```
✓ [82] cdf = pd.DataFrame(lm.coef_,X.columns,columns=['coeff'])
       cdf
```

```

✓ [82] coeff ⚭
area      2.180000e+02
bedrooms   5.093097e-14
bathrooms  6.560394e-14
parking    -9.063614e-14
stories_1   -3.530246e-14
stories_2   -8.039819e-15
stories_3   -1.754568e-14
stories_4   -4.830032e-14

✓ 0s ➔ predictions = lm.predict(X_test)
predictions

↳ array([1.0300000e+02, 2.0500000e+02, 1.8200000e+02, 6.4000000e+01,
       6.5000000e+01, 6.1000000e+01, 1.8900000e+02, 5.4000000e+01,
       7.6000000e+01, 1.2400000e+02, 8.1000000e+01, 2.1200000e+02,
       4.4000000e+01, 7.5000000e+01, 1.5300000e+02, 1.2200000e+02,
       4.0000000e+01, 1.6200000e+02, 5.6000000e+01, 9.3000000e+01,
       1.4600000e+02, 1.4500000e+02, 1.6300000e+02, 1.4500000e+02,
       9.7000000e+01, 9.0000000e+01, 3.2000000e+01, 8.1000000e+01,
       5.9000000e+01, 1.8000000e+02, 1.4500000e+02, 6.0000000e+00,
       1.9600000e+02, 4.0000000e+01, 1.3000000e+02, 4.8000000e+01,
       1.0700000e+02, 1.0600000e+02, 5.4000000e+01, 1.8000000e+01,
       9.8000000e+01, 2.1500000e+02, 1.1800000e+02, 8.1000000e+01,
       2.1000000e+01, 5.1000000e+01, 4.9000000e+01, 1.0200000e+02,
       1.5500000e+02, 8.1000000e+01, 6.5000000e+01, 1.2000000e+02,
       2.1700000e+02, 4.4000000e+01, 1.2100000e+02, 7.4000000e+01,
       4.7000000e+01, 1.1000000e+02, 9.4000000e+01, 5.9000000e+01,
       3.8000000e+01, 2.0000000e+02, 1.9700000e+02, 8.7000000e+01,
       1.0000000e+02, 1.5100000e+02, 4.6000000e+01, 1.7300000e+02,
       5.7000000e+01, 9.4000000e+01, 1.7500000e+02, 2.0900000e+02,
       8.3000000e+01, 4.8000000e+01, 4.0000000e+01, 1.0900000e+02,
       1.5000000e+01, 9.0000000e+00, 9.3000000e+01, 1.4200000e+02])

✓ 0s ➔ plt.scatter(y_test,predictions)
↳ <matplotlib.collections.PathCollection at 0x7f4cc01bfdf0>



```

```
✓ 1s  ➤ sns.distplot(y_test_predictions,bins=30)
↳ <ipython-input-85-5eafcd2b47918>:1: UserWarning:
    'distplot' is a deprecated function and will be removed in seaborn v0.14.0.
    Please adapt your code to use either 'displot' (a figure-level function with
    similar flexibility) or 'histplot' (an axes-level function for histograms).
```

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
<Axes: xlabel='price', ylabel='Density'>
```

